

# Project: Wrangling and Analyze Data

## Data Gathering

In the cell below, gather **all** three pieces of data for this project and load them in the notebook. **Note:** the methods required to gather each data are different.

1. Directly download the WeRateDogs Twitter archive data (twitter\_archive\_enhanced.csv)

```
In [1]: # import modules
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import shutil
import json
import tweepy
from tweepy import OAuthHandler
from timeit import default_timer as timer
from warnings import filterwarnings
import re
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

filterwarnings('ignore')
```

```
In [2]: # Load twitter archived dataset
twitter_archive = pd.read_csv('twitter-archive-enhanced.csv')
twitter_archive.head()
```


Out[2]:

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	timestamp	source	text	retweet
0	892420643555336193	NaN	NaN	2017-08-01 16:23:56 +0000	<a href="http://twitter.com/download/iphone" r...	This is Phineas. He's a mystical boy. Only eve...	
1	892177421306343426	NaN	NaN	2017-08-01 00:17:27 +0000	<a href="http://twitter.com/download/iphone" r...	This is Tilly. She's just checking pup on you....	
2	891815181378084864	NaN	NaN	2017-07-31 00:18:03 +0000	<a href="http://twitter.com/download/iphone" r...	This is Archie. He is a rare Norwegian Pouncin...	
3	891689557279858688	NaN	NaN	2017-07-30 15:58:51 +0000	<a href="http://twitter.com/download/iphone" r...	This is Darla. She commenced a snooze mid meal...	
4	891327558926688256	NaN	NaN	2017-07-29 16:00:24 +0000	<a href="http://twitter.com/download/iphone" r...	This is Franklin. He would like you to stop ca...	

In [3]: `twitter_archive.sample(10)`

Out[3]:

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	timestamp	source	
<b>1050</b>	743210557239623680	NaN	NaN	2016-06-15 22:36:19 +0000	<a href="http://twitter.com/download/iphone" r...	Meet Kay underground legend. P
<b>643</b>	793195938047070209	NaN	NaN	2016-10-31 21:00:23 +0000	<a href="http://twitter.com/download/iphone" r...	Say hello to Lily. pupset that her c
<b>1208</b>	715704790270025728	NaN	NaN	2016-04-01 00:58:13 +0000	<a href="http://vine.co" rel="nofollow">Vine -...	This is Bentley. He kisses back. 11/
<b>1769</b>	678380236862578688	NaN	NaN	2015-12-20 01:03:46 +0000	<a href="http://twitter.com/download/iphone" r...	This is Crump underestimat sno
<b>2310</b>	666786068205871104	NaN	NaN	2015-11-18 01:12:41 +0000	<a href="http://twitter.com/download/iphone" r...	Unfamiliar wi breed. Ears poi
<b>632</b>	793962221541933056	NaN	NaN	2016-11-02 23:45:19 +0000	<a href="http://twitter.com/download/iphone" r...	This is Maximu face is stuck like t
<b>1685</b>	681654059175129088	NaN	NaN	2015-12-29 01:52:46 +0000	<a href="http://twitter.com/download/iphone" r...	This is Toffee. happy pupper. Aq
<b>388</b>	826598365270007810	NaN	NaN	2017-02-01 01:09:42 +0000	<a href="http://twitter.com/download/iphone" r...	This is Pawnd... Pawnd. He's suz
<b>1659</b>	683098815881154561	NaN	NaN	2016-01-02 01:33:43 +0000	<a href="http://twitter.com/download/iphone" r...	aahhhhklsldhwnx 12/10 for be sn
<b>1550</b>	689154315265683456	NaN	NaN	2016-01-18 18:36:07 +0000	<a href="http://twitter.com/download/iphone" r...	We normally dor birds but I feel ba



## 1. Use the Requests library to download the tweet image prediction (image\_predictions.tsv)

```
In [4]: import requests
r = requests.get('https://d17h27t6h515a5.cloudfront.net/topher/2017/August/599fd2ad_image-predictions/image-predictions.tsv')
```

```
In [5]: # save image prediction data
if r.status_code == 200:
    with open('image_predictions.tsv', 'wb') as f:
        r.raw.decode_content = True
        shutil.copyfileobj(r.raw, f)
```

```
In [6]: image_prediction = pd.read_csv('image_predictions.tsv', sep='\t')
image_prediction.head()
```

Out[6]:

	tweet_id	jpg_url	img_num	p1	p1_conf	p1_dog	
0	666020888022790149	https://pbs.twimg.com/media/CT4udn0WwAA0aMy.jpg	1	Welsh_springer_spaniel	0.465074	True	
1	666029285002620928	https://pbs.twimg.com/media/CT42GRgUYAA5iDo.jpg	1	redbone	0.506826	True	miniature_p
2	666033412701032449	https://pbs.twimg.com/media/CT4521TWwAEvMyu.jpg	1	German_shepherd	0.596461	True	r
3	666044226329800704	https://pbs.twimg.com/media/CT5Dr8HUEAA-lEu.jpg	1	Rhodesian_ridgeback	0.408143	True	r
4	666049248165822465	https://pbs.twimg.com/media/CT5IQmsXIAAKY4A.jpg	1	miniature_pinscher	0.560311	True	Rc

In [7]: `image_prediction.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null int64
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
p2            2075 non-null object
p2_conf       2075 non-null float64
p2_dog        2075 non-null bool
p3            2075 non-null object
p3_conf       2075 non-null float64
p3_dog        2075 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB
```

1. Use the Tweepy library to query additional data via the Twitter API (tweet\_json.txt)

```
In [8]: # Setup tweepy api
consumer_key = 'None'
consumer_secret = 'None'
access_token = 'None'
access_secret = 'None'

auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
auth.set_access_token(access_token, access_secret)

api = tweepy.API(auth, wait_on_rate_limit=True)
```

```
In [ ]: # Extract tweets with Api
count = 0
fails_dict = {}
start = timer()
with open('tweet_json.txt', 'w') as f:
    for tweet_id in tweet_ids:
        count += 1
        print(str(count) + ": " + str(tweet_id))
        try:
            tweet = api.get_status(id=tweet_id, tweet_mode = 'extended')
            print('Success')
            json.dump(tweet._json, f)
            f.write('\n')
        except tweepy.TweepyException as e:
            print('Fail')
            fails_dict[tweet_id] = e
        pass

end = timer()
print(end - start)
print(fails_dict)
```

```
In [ ]: # Save tweet details in json format to a list of dictionaries
# read data line by line
with open('tweet_json.txt') as file:

    status = []
    for line in file:
        status.append(json.loads(line))
```

```
In [9]: tweets = pd.read_csv('tweets.csv')
        tweets.head()
```

Out[9]:

	tweet_id	retweet_count	favorite_count
0	892420643555336193	7009	33812
1	892177421306343426	5301	29329
2	891815181378084864	3481	22048
3	891689557279858688	7217	36899
4	891327558926688256	7760	35310

```
In [10]: tweets.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2327 entries, 0 to 2326
Data columns (total 3 columns):
tweet_id      2327 non-null int64
retweet_count 2327 non-null int64
favorite_count 2327 non-null int64
dtypes: int64(3)
memory usage: 54.6 KB
```



## Assessing Data

In this section, detect and document at least **eight (8) quality issues** and **two (2) tidiness issue**. You must use **both** visual assessment programmatic assesment to assess the data.

**Note:** pay attention to the following key points when you access the data.

- You only want original ratings (no retweets) that have images. Though there are 5000+ tweets in the dataset, not all are dog ratings and some are retweets.
- Assessing and cleaning the entire dataset completely would require a lot of time, and is not necessary to practice and demonstrate your skills in data wrangling. Therefore, the requirements of this project are only to assess and clean at least 8 quality issues and at least 2 tidiness issues in this dataset.
- The fact that the rating numerators are greater than the denominators does not need to be cleaned. This [unique rating system](http://knowyourmeme.com/memes/theyre-good-dogs-brent) (<http://knowyourmeme.com/memes/theyre-good-dogs-brent>) is a big part of the popularity of WeRateDogs.
- You do not need to gather the tweets beyond August 1st, 2017. You can, but note that you won't be able to gather the image predictions for these tweets since you don't have access to the algorithm used.

## Assessing twitter\_archive\_enhanced dataset

```
In [11]: twitter_archive.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id   78 non-null float64
in_reply_to_user_id     78 non-null float64
timestamp               2356 non-null object
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id     181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls           2297 non-null object
rating_numerator        2356 non-null int64
rating_denominator      2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
dtypes: float64(4), int64(3), object(10)
memory usage: 313.0+ KB
```

- retweet\_status\_id is float and not an int
- timestamp and retweeted\_status\_timestamp is object instead of datetime
- missing data in ( in\_reply\_to\_status\_id,in\_reply\_to\_user\_id ) and [retweet\_status\_id - expanding\_urls] )
- Duplicate tweets as retweet
- Nulls represented as None in name, doggo, floofer, pupper amd puppo
- tweets not related to dog ratings
- retweet\_status\_user\_id is float and not an int

```
In [12]: #twitter_archive[~twitter_archive.retweeted_status_id.isna()]['retweeted_status_id']
twitter_archive[twitter_archive.tweet_id == 888202515573088257]
```

Out[12]:

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	timestamp	source	text	retwee
19	888202515573088257	NaN	NaN	2017-07-21 01:02:36 +0000	href="http://twitter.com/download/iphone"	RT @dog_rates: This is Canela. She attempted s...	

Here we view a sample retweet. We can see that the data in the `text` column begins with **RT**. Also the `retweeted_status_id` has a value

```
In [13]: twitter_archive.describe()
```

Out[13]:

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	retweeted_status_id	retweeted_status_user_id	rating_numerator	rating_d
count	2.356000e+03	7.800000e+01	7.800000e+01	1.810000e+02	1.810000e+02	2356.000000	
mean	7.427716e+17	7.455079e+17	2.014171e+16	7.720400e+17	1.241698e+16	13.126486	
std	6.856705e+16	7.582492e+16	1.252797e+17	6.236928e+16	9.599254e+16	45.876648	
min	6.660209e+17	6.658147e+17	1.185634e+07	6.661041e+17	7.832140e+05	0.000000	
25%	6.783989e+17	6.757419e+17	3.086374e+08	7.186315e+17	4.196984e+09	10.000000	
50%	7.196279e+17	7.038708e+17	4.196984e+09	7.804657e+17	4.196984e+09	11.000000	
75%	7.993373e+17	8.257804e+17	4.196984e+09	8.203146e+17	4.196984e+09	12.000000	
max	8.924206e+17	8.862664e+17	8.405479e+17	8.874740e+17	7.874618e+17	1776.000000	

- minimum `rating_denominator` and `rating_numerator` is 0
- maximum `rating_denominator` and `rating_numerator` is 170 and 1776 respectively
- 75% of ratings value 1.2 or below

```
In [14]: twitter_archive[twitter_archive['rating_denominator'] == twitter_archive.rating_denominator.min()]
```

Out[14]:

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	timestamp	source	text	retv
313	835246439529840640	8.352460e+17	26259576.0	2017-02-24 21:54:03 +0000	href="http://twitter.com/download/iphone"	@jonnysun @Lin_Manuel ok jomny I know you're e...	

- One record has a rating denominator of 0

```
In [15]: twitter_archive[twitter_archive['rating_numerator'] == twitter_archive.rating_numerator.min()]
```

Out[15]:

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	timestamp	source	text	retwe
315	835152434251116546	NaN	NaN	2017-02-24 15:40:31 +0000	href="http://twitter.com/download/iphone"	When you're so blinded by your systematic plag...	
1016	746906459439529985	7.468859e+17	4.196984e+09	2016-06-26 03:22:31 +0000	href="http://twitter.com/download/iphone"	PUPDATE: can't see any. Even if I could, I cou...	

- Two records have a rating numerator or 0

## Assessing the tweets dataset

In [16]: `tweets.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2327 entries, 0 to 2326
Data columns (total 3 columns):
tweet_id      2327 non-null int64
retweet_count  2327 non-null int64
favorite_count 2327 non-null int64
dtypes: int64(3)
memory usage: 54.6 KB
```

- There are 2327 records for all the columns in the tweets dataset
- All records are of integer data types

In [17]: `tweets.describe().T`

Out[17]:

	count	mean	std	min	25%	50%	75%	max
<b>tweet_id</b>	2327.0	7.417930e+17	6.820795e+16	6.660209e+17	6.781394e+17	7.178418e+17	7.986547e+17	8.924206e+17
<b>retweet_count</b>	2327.0	2.471014e+03	4.184799e+03	1.000000e+00	4.945000e+02	1.148000e+03	2.860500e+03	7.073300e+04
<b>favorite_count</b>	2327.0	7.053779e+03	1.096152e+04	0.000000e+00	1.225000e+03	3.051000e+03	8.603500e+03	1.448750e+05

- mean retweet\_count is 2471 with the median of 1148. This indicates a heavily positively skewed data.
- mean favorite\_count is 7053 with median of 3051. This also indicates a positively skewed data
- minimum favorite and retweet counts are 0 and 1 respectively
- maximum favorite and retweet counts are 144875 and 70733 respectively

In [18]: `tweets.duplicated().sum()`

Out[18]: 0

In [19]: `tweets.tweet_id.nunique()`

Out[19]: 2327

- There are no duplicated rows in the data set

In [20]: `tweets.sample(10)`

Out[20]:

	tweet_id	retweet_count	favorite_count
<b>1223</b>	710997087345876993	1267	4223
<b>1297</b>	706166467411222528	1459	4742
<b>1366</b>	700002074055016451	1206	3038
<b>2101</b>	670093938074779648	277	911
<b>385</b>	824297048279236611	3569	14232
<b>1844</b>	675145476954566656	783	1918
<b>1696</b>	680085611152338944	8156	11762
<b>1578</b>	685667379192414208	511	2186
<b>1828</b>	675517828909424640	407	1151
<b>2197</b>	668274247790391296	198	721

- Sample of 10 records in the tweets dataset.

## Assessing the image\_prediction dataset

In [21]: `image_prediction.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null int64
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
p2            2075 non-null object
p2_conf       2075 non-null float64
p2_dog        2075 non-null bool
p3            2075 non-null object
p3_conf       2075 non-null float64
p3_dog        2075 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB
```

- There are 2075 observations in the data set. The data has four distinct data types (int, float, bool and object)

In [22]: `image_prediction.describe()`

Out[22]:

	tweet_id	img_num	p1_conf	p2_conf	p3_conf
<b>count</b>	2.075000e+03	2075.000000	2075.000000	2.075000e+03	2.075000e+03
<b>mean</b>	7.384514e+17	1.203855	0.594548	1.345886e-01	6.032417e-02
<b>std</b>	6.785203e+16	0.561875	0.271174	1.006657e-01	5.090593e-02
<b>min</b>	6.660209e+17	1.000000	0.044333	1.011300e-08	1.740170e-10
<b>25%</b>	6.764835e+17	1.000000	0.364412	5.388625e-02	1.622240e-02
<b>50%</b>	7.119988e+17	1.000000	0.588230	1.181810e-01	4.944380e-02
<b>75%</b>	7.932034e+17	1.000000	0.843855	1.955655e-01	9.180755e-02
<b>max</b>	8.924206e+17	4.000000	1.000000	4.880140e-01	2.734190e-01

- The mean first prediction confidence is ~0.59.
- 25% of the time, the algorithm predicts with a confidence of over 0.84

In [23]: `image_prediction.tweet_id.duplicated().sum()`

Out[23]: 0

In [24]: `image_prediction.tweet_id.nunique()`

Out[24]: 2075

In [25]: `image_prediction.sample(10)`

Out[25]:

	tweet_id	jpg_url	img_num	p1	p1_conf	p1_dog	
1742	822647212903690241	https://pbs.twimg.com/media/C2oRbOuWEAAbVSI.jpg	1	Samoyed	0.416769	True	
673	683357973142474752	https://pbs.twimg.com/media/CXvGbWeWMAcRbyJ.jpg	1	Pembroke	0.406509	True	
1732	821149554670182400	https://pbs.twimg.com/ext_tw_video_thumb/82114...	1	German_shepherd	0.515933	True	
1748	823581115634085888	https://pbs.twimg.com/media/C23ypm6VQAAO31l.jpg	1	dingo	0.280949	False	(
2040	885167619883638784	https://pbs.twimg.com/media/DEi_N9qXYAAgEEw.jpg	4	malamute	0.812482	True	
1071	716439118184652801	https://pbs.twimg.com/media/CfFNk7cWAAA-hND.jpg	1	Siberian_husky	0.396495	True	
1482	781163403222056960	https://pbs.twimg.com/media/Ctc_-BTWEAAQpZh.jpg	1	Shetland_sheepdog	0.973841	True	
1502	784517518371221505	https://pbs.twimg.com/media/CuMqhGrXYAQwRqU.jpg	2	malamute	0.757764	True	
150	668641109086707712	https://pbs.twimg.com/media/CUD9ivxWUAAuXSQ.jpg	1	vacuum	0.432594	False	
1170	736225175608430592	https://pbs.twimg.com/media/CjeY5DKXEEA3WkD.jpg	1	Labrador_retriever	0.399217	True	West_High

- There are no duplicated rows in the data set
- All records are unique



## Quality issues

### *twitter\_archive\_enhanced dataset*

1. retweet\_status\_id and retweet\_status\_user\_id is float and not an int
2. timestamp is object instead of datetime
3. missing data in ( in\_reply\_to\_status\_id,in\_reply\_to\_user\_id )
4. Duplicate tweets as retweet
5. Invalid Dog Names
6. incorrect rating denominators
7. incorrect rating numerators with decimal values incorrectly extracted

### *image\_prediction dataset*

1. Some dog names in the prediction columns P1, P2 and P3 are not capitalised

## Tidiness issues

1. dog stages represented as separate columns
2. The `tweets` table should be a part of the `twitter_archive_enhanced` table

## Cleaning Data

In this section, clean **all** of the issues you documented while assessing.

**Note:** Make a copy of the original data before cleaning. Cleaning includes merging individual pieces of data according to the rules of [tidy data](https://cran.r-project.org/web/packages/tidyr/vignettes/tidy-data.html) (<https://cran.r-project.org/web/packages/tidyr/vignettes/tidy-data.html>). The result should be a high-quality and tidy master pandas DataFrame (or DataFrames, if appropriate).

```
In [26]: # Make copies of original pieces of data

twitter_archive_clean = twitter_archive.copy()
tweets_clean = tweets.copy()
image_prediction_clean = image_prediction.copy()
```

# Cleaning

## Quality Issues

### Issue #1

#### Define:

Convert retweet\_status\_id and retweet\_status\_user\_id from float to string.

#### Code

```
In [27]: twitter_archive_clean.retweeted_status_id = twitter_archive_clean.retweeted_status_id.astype(str)
twitter_archive_clean.retweeted_status_user_id = twitter_archive_clean.retweeted_status_user_id.astype(str)
```

#### Test

```
In [28]: twitter_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id   78 non-null float64
in_reply_to_user_id     78 non-null float64
timestamp               2356 non-null object
source                 2356 non-null object
text                   2356 non-null object
retweeted_status_id     2356 non-null object
retweeted_status_user_id 2356 non-null object
retweeted_status_timestamp 181 non-null object
expanded_urls          2297 non-null object
rating_numerator        2356 non-null int64
rating_denominator      2356 non-null int64
name                   2356 non-null object
doggo                  2356 non-null object
floofer                2356 non-null object
pupper                2356 non-null object
puppo                  2356 non-null object
dtypes: float64(2), int64(3), object(12)
memory usage: 313.0+ KB
```

retweet\_status\_id and retweet\_status\_user\_id now have **object** as the datatype

## Issue #2

### Define:

Convert the data type of timestamp from object to datetime

### Code

```
In [29]: twitter_archive_clean.timestamp = pd.to_datetime(twitter_archive_clean.timestamp)
```

## Test

```
In [30]: twitter_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id    78 non-null float64
in_reply_to_user_id      78 non-null float64
timestamp               2356 non-null datetime64[ns]
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id      2356 non-null object
retweeted_status_user_id 2356 non-null object
retweeted_status_timestamp 181 non-null object
expanded_urls            2297 non-null object
rating_numerator          2356 non-null int64
rating_denominator        2356 non-null int64
name                     2356 non-null object
doggo                    2356 non-null object
floofer                  2356 non-null object
pupper                   2356 non-null object
puppo                    2356 non-null object
dtypes: datetime64[ns](1), float64(2), int64(3), object(11)
memory usage: 313.0+ KB
```

timestamp column is now of **datetime** data type

## Issue #3

**Define:**

drop the in\_reply\_to\_status\_id and in\_reply\_to\_user\_id columns

**Code**

```
In [31]: twitter_archive_clean.drop(['in_reply_to_status_id', 'in_reply_to_user_id'], axis=1, inplace=True)
```

**Test**

```
In [32]: twitter_archive_clean.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 15 columns):
tweet_id                2356 non-null int64
timestamp               2356 non-null datetime64[ns]
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id     2356 non-null object
retweeted_status_user_id 2356 non-null object
retweeted_status_timestamp 181 non-null object
expanded_urls           2297 non-null object
rating_numerator        2356 non-null int64
rating_denominator      2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
dtypes: datetime64[ns](1), int64(3), object(11)
memory usage: 276.2+ KB
```

in\_reply\_to\_status\_id and in\_reply\_to\_user\_id columns are no more in the dataset

## Issue #4

### Define :

Create a dataframe of retweets and drop all rows where retweet\_status\_id is present.

### Code

```
In [33]: # create a data frame of retweets
retweets = twitter_archive_clean[~twitter_archive_clean.retweeted_status_timestamp.isna()]
```

```
In [34]: retweets.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 181 entries, 19 to 2260
Data columns (total 15 columns):
tweet_id                181 non-null int64
timestamp               181 non-null datetime64[ns]
source                  181 non-null object
text                    181 non-null object
retweeted_status_id     181 non-null object
retweeted_status_user_id 181 non-null object
retweeted_status_timestamp 181 non-null object
expanded_urls           180 non-null object
rating_numerator        181 non-null int64
rating_denominator      181 non-null int64
name                    181 non-null object
doggo                   181 non-null object
floofer                 181 non-null object
pupper                  181 non-null object
puppo                   181 non-null object
dtypes: datetime64[ns](1), int64(3), object(11)
memory usage: 22.6+ KB
```

```
In [35]: # delete retweets from the twitter_archive_clean dataset
twitter_archive_clean.drop(retweets.index, axis=0, inplace=True)
```

```
In [36]: # drop all retweet related columns since there are no records for them
twitter_archive_clean.drop(['retweeted_status_id', 'retweeted_status_user_id', 'retweeted_status_timestamp'], axis=1, inplace=True)
```

## Test

```
In [37]: twitter_archive_clean.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2175 entries, 0 to 2355
Data columns (total 12 columns):
tweet_id          2175 non-null int64
timestamp         2175 non-null datetime64[ns]
source            2175 non-null object
text              2175 non-null object
expanded_urls     2117 non-null object
rating_numerator  2175 non-null int64
rating_denominator 2175 non-null int64
name              2175 non-null object
doggo             2175 non-null object
floofer           2175 non-null object
pupper            2175 non-null object
puppo             2175 non-null object
dtypes: datetime64[ns](1), int64(3), object(8)
memory usage: 220.9+ KB
```

- There are no records for retweets in the dataset. The data have only unique tweets now
- All retweet related columns have been removed.

## Issue #5

**Define:**

- Isolate all rows that have invalid dog names.
- Check if we can further extract some of the names
- Update list of dog names
- Clean up the dog names to only contain valid names

**Code**

```
In [38]: mask = twitter_archive_clean.name.str.contains('^[a-z]', regex=True)
twitter_archive_clean[mask][['text', 'name']].head()
```

Out[38]:

	text	name
22	I've yet to rate a Venezuelan Hover Wiener. Th...	such
56	Here is a pupper approaching maximum borkdrive...	a
169	We only rate dogs. This is quite clearly a smo...	quite
193	Guys, we only rate dogs. This is quite clearly...	quite
335	There's going to be a dog terminal at JFK Airp...	not

- There are 104 observations.
- There are some names that were not extracted properly from the text. We will try and extract some more names

```
In [39]: dog_names_clean = twitter_archive_clean[mask][['text', 'name']]
```



```
In [40]: # Create a list
# we will use the list to store the extracted dog_names from our execption list
dg_name = []
for dog_name in dog_names_clean['text']:
    try:
        name = re.search('(<=named)(.\w+)', dog_name)
        if name:
            name = name.group(1)
        else:
            name = re.search('(<=name is)(.\w+)', dog_name)
            name = name.group(1)

    except:
        pass
    dg_name.append(name)
```

```
In [41]: dog_name_fix = pd.Series(dg_name,dog_names_clean.index)
```

```
In [42]: twitter_archive_clean['names_fixed'] = dog_name_fix
```

```
In [43]: twitter_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2175 entries, 0 to 2355
Data columns (total 13 columns):
tweet_id          2175 non-null int64
timestamp         2175 non-null datetime64[ns]
source            2175 non-null object
text              2175 non-null object
expanded_urls     2117 non-null object
rating_numerator  2175 non-null int64
rating_denominator 2175 non-null int64
name              2175 non-null object
doggo             2175 non-null object
floofer           2175 non-null object
pupper            2175 non-null object
puppo             2175 non-null object
names_fixed       22 non-null object
dtypes: datetime64[ns](1), int64(3), object(9)
memory usage: 237.9+ KB
```

- There are 22 new dog names that we have been able to fetch

```
In [44]: # UPDATE THE DOG NAMES
twitter_archive_clean.names_fixed.fillna(twitter_archive_clean.name,inplace=True)
```

## Test

```
In [45]: mask = twitter_archive_clean.names_fixed.str.contains('^[a-z]', regex=True)
twitter_archive_clean[mask]['names_fixed'].value_counts()
```

```
Out[45]: a                35
the                 8
an                  5
very                4
one                 4
just                3
quite               3
not                 2
actually            2
getting             2
mad                 1
old                 1
life                1
infuriating         1
officially           1
all                  1
such                 1
unacceptable         1
his                  1
by                   1
light                1
space                1
this                 1
incredibly           1
Name: names_fixed, dtype: int64
```

- The number of wrong dog names have reduce

```
In [46]: twitter_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 2175 entries, 0 to 2355  
Data columns (total 13 columns):  
tweet_id          2175 non-null int64  
timestamp         2175 non-null datetime64[ns]  
source           2175 non-null object  
text             2175 non-null object  
expanded_urls    2117 non-null object  
rating_numerator 2175 non-null int64  
rating_denominator 2175 non-null int64  
name             2175 non-null object  
doggo            2175 non-null object  
floofer         2175 non-null object  
pupper          2175 non-null object  
puppo           2175 non-null object  
names_fixed      2175 non-null object  
dtypes: datetime64[ns](1), int64(3), object(9)  
memory usage: 237.9+ KB
```

```
In [47]: # remove leading spaces from names_fixed  
twitter_archive_clean.names_fixed = twitter_archive_clean.names_fixed.str.replace(' ', '')
```

```
In [48]: twitter_archive_clean[twitter_archive_clean.name == 'Zoey']['name'].value_counts()
```

```
Out[48]: Zoey      3  
         Name: name, dtype: int64
```

```
In [49]: twitter_archive_clean[twitter_archive_clean.names_fixed == 'Zoey']['names_fixed'].value_counts()
```

```
Out[49]: Zoey      4  
         Name: names_fixed, dtype: int64
```

- names\_fixed column we can confirm has more names than the original names column. We will proceed to replace the names column with the names\_fixed column

```
In [50]: # drop old column name  
twitter_archive_clean.drop('name',axis=1, inplace=True)
```

```
In [51]: # Rename column name  
twitter_archive_clean.rename(columns={'names_fixed':'name'}, inplace=True)
```

```
In [52]: twitter_archive_clean.info()  
  
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 2175 entries, 0 to 2355  
Data columns (total 12 columns):  
tweet_id          2175 non-null int64  
timestamp         2175 non-null datetime64[ns]  
source            2175 non-null object  
text              2175 non-null object  
expanded_urls     2117 non-null object  
rating_numerator  2175 non-null int64  
rating_denominator 2175 non-null int64  
doggo             2175 non-null object  
floofer           2175 non-null object  
pupper            2175 non-null object  
puppo             2175 non-null object  
name              2175 non-null object  
dtypes: datetime64[ns](1), int64(3), object(8)  
memory usage: 220.9+ KB
```

## Issue #6

### Define:

- Update incorrect ratings numerator and denominators

## Code

```
In [53]: #create a dataframe of all denominators greater than 10
denom_to_fix = twitter_archive_clean[twitter_archive_clean.rating_denominator != 10][['text', 'rating_denominator']]

In [54]: # Extract denominator from tweet text
rating_denom = []

for text in denom_to_fix['text']:
    split_text = text.split('https')
    word = split_text[0]

    rating_denom.append(re.search('(\d+)(?!.*\d)', word).group(1))

In [55]: # Convert ratings_denom to pandas series and append to denom_to_fix
rating_denom = pd.Series(rating_denom, denom_to_fix.index)

In [56]: # convert datatypes from string and int to float
rating_denom = rating_denom.astype(float)
twitter_archive_clean.rating_denominator = twitter_archive_clean.rating_denominator.astype(float)

In [57]: # add rating_denominator to twitter_archive_clean dataframe
twitter_archive_clean['rating_denom'] = rating_denom

In [58]: # Update rating_denom with values from rating_denominator
twitter_archive_clean['rating_denom'] = twitter_archive_clean.rating_denom.fillna(twitter_archive_clean.rating_denominator)
```

## Test

```
In [59]: twitter_archive_clean[twitter_archive_clean.rating_denom != twitter_archive_clean.rating_denominator]
```

```
Out[59]:
```

	tweet_id	timestamp	source	text	expanded_urls
313	835246439529840640	2017-02-24 21:54:03	<a href="http://twitter.com/download/iphone" r...	@jonnysun @Lin_Manuel ok jomny I know you're e...	NaN
1068	740373189193256964	2016-06-08 02:41:38	<a href="http://twitter.com/download/iphone" r...	After so many requests, this is Bretagne. She ...	https://twitter.com/dog_rates/status/740373189...
1165	722974582966214656	2016-04-21 02:25:47	<a href="http://twitter.com/download/iphone" r...	Happy 4/20 from the squad! 13/10 for all https...	https://twitter.com/dog_rates/status/722974582...
1202	716439118184652801	2016-04-03 01:36:11	<a href="http://twitter.com/download/iphone" r...	This is Bluebert. He just saw that both #Final...	https://twitter.com/dog_rates/status/716439118...
1662	682962037429899265	2016-01-01 16:30:13	<a href="http://twitter.com/download/iphone" r...	This is Darrel. He just robbed a 7/11 and is i...	https://twitter.com/dog_rates/status/682962037...
2335	666287406224695296	2015-11-16 16:11:11	<a href="http://twitter.com/download/iphone" r...	This is an Albanian 3 1/2 legged Episcopalian...	https://twitter.com/dog_rates/status/666287406...

```
In [60]: twitter_archive_clean.drop('rating_denominator', axis=1, inplace=True)
```

```
In [61]: twitter_archive_clean.rename(columns={'rating_denom': 'rating_denominator'}, inplace=True)
```

```
In [62]: twitter_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2175 entries, 0 to 2355
Data columns (total 12 columns):
tweet_id          2175 non-null int64
timestamp         2175 non-null datetime64[ns]
source            2175 non-null object
text              2175 non-null object
expanded_urls     2117 non-null object
rating_numerator  2175 non-null int64
doggo             2175 non-null object
floofer           2175 non-null object
pupper           2175 non-null object
puppo            2175 non-null object
name              2175 non-null object
rating_denominator 2175 non-null float64
dtypes: datetime64[ns](1), float64(1), int64(2), object(8)
memory usage: 220.9+ KB
```

- All incorrect `rating_denominator` issues have been resolved

## Issue #7

### Define:

- capture ratings from text column in the dataset
- extract numerator values from text
- update numerator values in the dataset
- convert data type for `rating_numerator` to float

### Code

```
In [63]: # Extract ratings info from text
ratings = twitter_archive_clean.text.str.extract('((?:\d+\.)?\d+)\./(\d+)', expand=True)
```

```
In [64]: ratings.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2175 entries, 0 to 2355
Data columns (total 2 columns):
0      2175 non-null object
1      2175 non-null object
dtypes: object(2)
memory usage: 51.0+ KB
```

```
In [65]: #update rating numerator values in dataset
twitter_archive_clean.rating_numerator = ratings[0]
```

```
In [66]: # convert rating_numerator to float
twitter_archive_clean.rating_numerator = twitter_archive_clean.rating_numerator.astype(float)
```

## Test



In [67]: `twitter_archive_clean.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2175 entries, 0 to 2355
Data columns (total 12 columns):
tweet_id          2175 non-null int64
timestamp         2175 non-null datetime64[ns]
source            2175 non-null object
text              2175 non-null object
expanded_urls     2117 non-null object
rating_numerator  2175 non-null float64
doggo             2175 non-null object
floofer           2175 non-null object
pupper           2175 non-null object
puppo             2175 non-null object
name              2175 non-null object
rating_denominator 2175 non-null float64
dtypes: datetime64[ns](1), float64(2), int64(1), object(8)
memory usage: 220.9+ KB
```

In [68]: `twitter_archive_clean.rating_numerator.unique()`

```
Out[68]: array([ 1.30000000e+01,  1.20000000e+01,  1.40000000e+01,
 1.35000000e+01,  1.70000000e+01,  1.10000000e+01,
 1.00000000e+01,  4.20000000e+02,  6.66000000e+02,
 6.00000000e+00,  1.82000000e+02,  1.50000000e+01,
 9.60000000e+02,  0.00000000e+00,  7.00000000e+00,
 8.40000000e+01,  2.40000000e+01,  9.75000000e+00,
 5.00000000e+00,  1.12700000e+01,  3.00000000e+00,
 8.00000000e+00,  9.00000000e+00,  4.00000000e+00,
 1.65000000e+02,  1.77600000e+03,  2.04000000e+02,
 5.00000000e+01,  9.90000000e+01,  8.00000000e+01,
 4.50000000e+01,  6.00000000e+01,  4.40000000e+01,
 1.00000000e+00,  1.43000000e+02,  1.21000000e+02,
 2.00000000e+01,  9.50000000e+00,  1.12600000e+01,
 2.00000000e+00,  1.44000000e+02,  8.80000000e+01])
```

- `rating_numerator` datatype is now float

## Issue #8

### Define:

- Capitalise first letter in dog names in the prediction columns P1, P2 and P3

### Code

```
In [69]: image_prediction_clean.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null int64
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
p2            2075 non-null object
p2_conf       2075 non-null float64
p2_dog        2075 non-null bool
p3            2075 non-null object
p3_conf       2075 non-null float64
p3_dog        2075 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB
```

```
In [70]: # confirm names not capitalised in the column
capital_check = image_prediction_clean.p1.str.contains('^[a-z]', regex=True)
image_prediction_clean[capital_check]['p1'].head()
```

```
Out[70]: 1          redbone
4    miniature_pinscher
6          box_turtle
7            chow
8      shopping_cart
Name: p1, dtype: object
```

```
In [71]: image_prediction_clean.p1 = image_prediction_clean.p1.str.title()
image_prediction_clean.p2 = image_prediction_clean.p2.str.title()
image_prediction_clean.p3 = image_prediction_clean.p3.str.title()
```

## Test

```
In [72]: capital_check = image_prediction_clean.p1.str.contains('^[a-z]', regex=True)
image_prediction_clean[capital_check]['p1'].head()
```

```
Out[72]: Series([], Name: p1, dtype: object)
```

```
In [73]: capital_check = image_prediction_clean.p2.str.contains('^[a-z]', regex=True)
image_prediction_clean[capital_check]['p2'].head()
```

```
Out[73]: Series([], Name: p2, dtype: object)
```

```
In [74]: capital_check = image_prediction_clean.p3.str.contains('^[a-z]', regex=True)
image_prediction_clean[capital_check]['p3'].head()
```

```
Out[74]: Series([], Name: p3, dtype: object)
```

- P1, P2 and P3 columns all have first letters capitalised

# Tidiness

## Issue #1

### Define:

- create a dog\_stage column
- Merge all dog stage columns together
- drop individual doggo, puppo, floofer, pupper columns

### Code

```
In [75]: # extract the various dog_stage columns
doggo = twitter_archive_clean[twitter_archive_clean.doggo == 'doggo']['doggo']
floofer = twitter_archive_clean[twitter_archive_clean.floofer == 'floofer']['floofer']
pupper = twitter_archive_clean[twitter_archive_clean.pupper == 'pupper']['pupper']
puppo = twitter_archive_clean[twitter_archive_clean.puppo == 'puppo']['puppo']
```

```
In [76]: # merge all dog_stage into one column
dog_stage = doggo.append(floofer)
dog_stage = dog_stage.append(pupper)
dog_stage = dog_stage.append(puppo)
```

```
In [77]: dg_index, dg_value = list(dog_stage.index), list(dog_stage.values)
```

```
In [78]: # create an empty dog_stage column
twitter_archive_clean['dog_stage'] = None
```

```
In [79]: # populate dog_stage column
for i in range(len(dg_value)):
    twitter_archive_clean.dog_stage.loc[dg_index[i]] = dg_value[i]
```

```
In [80]: # drop doggo, puppo, floofer and pupper columns. This has already been merged as dog_stage
twitter_archive_clean.drop(['doggo', 'puppo', 'floofer', 'pupper'], axis=1, inplace=True)
```

## Test

```
In [81]: twitter_archive_clean.dog_stage.value_counts()
```

```
Out[81]: pupper      234
doggo      75
puppo      25
floofer    10
Name: dog_stage, dtype: int64
```

```
In [82]: twitter_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2175 entries, 0 to 2355
Data columns (total 9 columns):
tweet_id      2175 non-null int64
timestamp     2175 non-null datetime64[ns]
source        2175 non-null object
text          2175 non-null object
expanded_urls 2117 non-null object
rating_numerator 2175 non-null float64
name          2175 non-null object
rating_denominator 2175 non-null float64
dog_stage     344 non-null object
dtypes: datetime64[ns](1), float64(2), int64(1), object(5)
memory usage: 249.9+ KB
```

- dog\_stage column has been added to the data set.
- the individual dog\_stages (pupper, floofer, doggo, puppo) have been merged into one column

```
In [ ]:
```

## Issue #2

### Define:

Add tweets table to twitter\_archive\_enhanced table

### Code

```
In [83]: twitter_archive_clean = twitter_archive_clean.merge(tweets_clean, how='left', on='tweet_id')
```

```
In [ ]:
```

### Test

```
In [84]: #confirm that the merge is successful  
twitter_archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 2175 entries, 0 to 2174  
Data columns (total 11 columns):  
tweet_id          2175 non-null int64  
timestamp         2175 non-null datetime64[ns]  
source            2175 non-null object  
text              2175 non-null object  
expanded_urls     2117 non-null object  
rating_numerator  2175 non-null float64  
name              2175 non-null object  
rating_denominator 2175 non-null float64  
dog_stage         344 non-null object  
retweet_count     2167 non-null float64  
favorite_count    2167 non-null float64  
dtypes: datetime64[ns](1), float64(4), int64(1), object(5)  
memory usage: 203.9+ KB
```

- Retweets and Favorite Counts have been added to the twitter\_archive dataset

## Storing Data

*Two datasets are to be saved.*

1. The twitter archived enhanced master dataset
2. The image\_prediction\_data set

```
In [85]: twitter_archive_clean.to_csv('twitter_archive_master.csv', index=False)  
         image_prediction_clean.to_csv('image_prediction.csv', index=False)
```

## Analyzing and Visualizing Data

In this section, analyze and visualize your wrangled data. You must produce at least **three (3) insights and one (1) visualization**.

### *Analyzing Twitter Archived Master Dataset*

```
In [86]: # read data and view first few rows
df = pd.read_csv('twitter_archive_master.csv')
df.head()
```

Out[86]:

	tweet_id	timestamp	source	text	expanded_urls	rating
0	892420643555336193	2017-08-01 16:23:56	<a href="http://twitter.com/download/iphone" r...	This is Phineas. He's a mystical boy. Only eve...	https://twitter.com/dog_rates/status/892420643...	
1	892177421306343426	2017-08-01 00:17:27	<a href="http://twitter.com/download/iphone" r...	This is Tilly. She's just checking pup on you....	https://twitter.com/dog_rates/status/892177421...	
2	891815181378084864	2017-07-31 00:18:03	<a href="http://twitter.com/download/iphone" r...	This is Archie. He is a rare Norwegian Pouncin...	https://twitter.com/dog_rates/status/891815181...	
3	891689557279858688	2017-07-30 15:58:51	<a href="http://twitter.com/download/iphone" r...	This is Darla. She commenced a snooze mid meal...	https://twitter.com/dog_rates/status/891689557...	
4	891327558926688256	2017-07-29 16:00:24	<a href="http://twitter.com/download/iphone" r...	This is Franklin. He would like you to stop ca...	https://twitter.com/dog_rates/status/891327558...	



In [ ]:



```
In [87]: # sample 10 records from the dataset  
df.sample(10)
```

Out[87]:

	tweet_id	timestamp	source	text	expanded_urls
933	733482008106668032	2016-05-20 02:18:32	<a href="http://twitter.com/download/iphone" r...	"Ello this is dog how may I assist" ...10/10 h...	https://twitter.com/dog_rates/status/733482008...
614	778383385161035776	2016-09-21 00:00:35	<a href="http://twitter.com/download/iphone" r...	This is Nala. She's a future Dogue model. Won'...	https://twitter.com/dog_rates/status/778383385...
2156	666268910803644416	2015-11-16 14:57:41	<a href="http://twitter.com/download/iphone" r...	Very concerned about fellow dog trapped in com...	https://twitter.com/dog_rates/status/666268910...
508	793962221541933056	2016-11-02 23:45:19	<a href="http://twitter.com/download/iphone" r...	This is Maximus. His face is stuck like that. ...	https://twitter.com/dog_rates/status/793962221...
362	820749716845686786	2017-01-15 21:49:15	<a href="http://twitter.com/download/iphone" r...	Meet Sunny. He can take down a polar bear in o...	https://twitter.com/dog_rates/status/820749716...
325	826848821049180160	2017-02-01 17:44:55	<a href="http://twitter.com/download/iphone" r...	This is Cupid. He was found in the trash. Now ...	https://twitter.com/dog_rates/status/826848821...
187	850333567704068097	2017-04-07 13:04:55	<a href="http://twitter.com/download/iphone" r...	@markhoppus MARK THAT DOG HAS SEEN AND EXPERIE...	NaN
921	735991953473572864	2016-05-27 00:32:10	<a href="http://twitter.com/download/iphone" r...	This is Maxaroni. He's curly af. Also rather f...	https://twitter.com/dog_rates/status/735991953...
1392	687704180304273409	2016-01-14 18:33:48	<a href="http://twitter.com/download/iphone" r...	Say hello to Blakely. He thinks that's a hat. ...	https://twitter.com/dog_rates/status/687704180...

	tweet_id	timestamp	source	text	expanded_urls
775	752173152931807232	2016-07-10 16:10:29	<a href="http://twitter.com/download/iphone" r...	This is Brody. He's a lifeguard. Always prepar...	https://twitter.com/dog_rates/status/752173152...



In [88]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2175 entries, 0 to 2174
Data columns (total 11 columns):
tweet_id          2175 non-null int64
timestamp         2175 non-null object
source            2175 non-null object
text              2175 non-null object
expanded_urls     2117 non-null object
rating_numerator  2175 non-null float64
name              2175 non-null object
rating_denominator 2175 non-null float64
dog_stage         344 non-null object
retweet_count     2167 non-null float64
favorite_count    2167 non-null float64
dtypes: float64(4), int64(1), object(6)
memory usage: 187.0+ KB
```

- There are 2175 observations in the dataset.
- We only have 2167 records for retweets and favorites


In [89]: `df.describe()`

Out[89]:


	tweet_id	rating_numerator	rating_denominator	retweet_count	favorite_count
<b>count</b>	2.175000e+03	2175.000000	2175.000000	2167.000000	2167.000000
<b>mean</b>	7.371205e+17	13.177140	10.477241	2236.188740	7574.592986
<b>std</b>	6.748668e+16	47.705088	6.958764	3983.510947	11184.086552
<b>min</b>	6.660209e+17	0.000000	7.000000	1.000000	45.000000
<b>25%</b>	6.768432e+17	10.000000	10.000000	477.000000	1602.500000
<b>50%</b>	7.098528e+17	11.000000	10.000000	1058.000000	3409.000000
<b>75%</b>	7.894226e+17	12.000000	10.000000	2550.500000	9393.500000
<b>max</b>	8.924206e+17	1776.000000	170.000000	70733.000000	144875.000000

- Mean rating\_denominator is ~10 with mean rating\_numerator ~13.17
- 75% of the rating denominator have a value of 10 whereas that of rating numerator is at most 12

```
In [90]: df.plot(x='favorite_count', y='retweet_count', kind='scatter')  
plt.title('Plot of Retweets against Favorites')  
plt.show()
```




```
In [91]: sns.countplot(data = df, x='dog_stage')  
plt.title('Distribution of Dog Stages');
```




- Pupper is the most popular dog\_stage

```
In [92]: df.retweet_count.hist()  
plt.title('Histogram of Retweet Counts');
```



- The histogram of retweet counts is positively skewed.

```
In [93]: df.favorite_count.hist();  
plt.title('Histogram of Favorite Counts');
```




## Insights:

1. Most popular dog\_stage is Pupper
2. There is a positive correlation between favorites and retweets. If a tweet is retweeted, there is a high chance that it will be favorited as well
3. Favorite and Retweet counts are heavily skewed to the right. This means that only a small portion of tweets are highly retweeted or favorited

## Visualization

```
In [94]: plt.figure(figsize=(8,5))  
plt.scatter(x = df.retweet_count, y = df.favorite_count)  
plt.xlabel('Retweets')  
plt.ylabel('Favourite')  
plt.title('Relationship between Favorites and Retweets')  
plt.show()
```



```
In [ ]:
```